

Tensorflow



Agenda

- Tensorflow Introduction
- Conferences
- XLA
 - just-in-time (JIT)
 - ahead-of-time (AOT)
- JIT and AOT examples

Tensorflow

TensorFlow™ is an open source software library for numerical computation using data flow graphs. Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them.

Tensors

A tensor consists of a set of primitive values shaped into an array of any number of dimensions.

```
3 # a rank 0 tensor; this is a scalar with shape []  
[1., 2., 3.] # a rank 1 tensor; this is a vector with shape [3]  
[[1., 2., 3.], [4., 5., 6.]] # a rank 2 tensor; a matrix with shape [2, 3]  
[[[1., 2., 3.]], [[7., 8., 9.]]] # a rank 3 tensor with shape [2, 1, 3]
```

Rank and Shape

Rank	Math entity
0	Scalar (magnitude only)
1	Vector (magnitude and direction)
2	Matrix (table of numbers)
3	3-Tensor (cube of numbers)
n	n-Tensor (you get the idea)

Rank	Shape	Dimension number	Example
0	[]	0-D	A 0-D tensor. A scalar.
1	[D0]	1-D	A 1-D tensor with shape [5].
2	[D0, D1]	2-D	A 2-D tensor with shape [3, 4].
3	[D0, D1, D2]	3-D	A 3-D tensor with shape [1, 4, 3].
n	[D0, D1, ... Dn-1]	n-D	A tensor with shape [D0, D1, ... Dn-1].

Constants

Constants Variables have a value that does not change during runtime.

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

Placeholders

A graph can be parameterized to accept external inputs, known as placeholders. A placeholder is a promise to provide a value later.

```
a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
adder_node = a + b  # + provides a shortcut for tf.add(a, b)
```

```
print(sess.run(adder_node, {a: 3, b: 4.5}))
print(sess.run(adder_node, {a: [1, 3], b: [2, 4]}))
```

Variables

To make the model trainable, we need to be able to modify the graph to get new outputs with the same input. Variables allow us to add trainable parameters to a graph.


```
W = tf.Variable([.3], dtype=tf.float32)
b = tf.Variable([-.3], dtype=tf.float32)
x = tf.placeholder(tf.float32)
linear_model = W * x + b
sess = tf.Session()
print(sess.run(linear_model, {x: [1, 2, 3, 4]}))
```


Computational Graph

A computational graph is a series of TensorFlow operations arranged into a graph of nodes

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
node3 = tf.add(node1, node2)
sess = tf.Session()
print("sess.run(node3):", sess.run(node3))

//Print result: sess.run(node3): 7.0
```



Session


A **session** encapsulates the control and state of the TensorFlow runtime.

```
# Create a default in-process session.  
with tf.Session() as sess:  
    # ...  
  
# Create a remote session.  
with tf.Session("grpc://example.org:2222"):  
    # ...
```

Computational Graph

```
a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
adder_node = a + b  # + provides a shortcut for tf.add(a, b)
sess = tf.Session()
print(sess.run(adder_node, {a: 3, b: 4.5}))
print(sess.run(adder_node, {a: [1, 3], b: [2, 4]}))


//print results:
7.5
[ 3.  7.]
```



Computational Graph

```
a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
adder_node = a + b  # + provides a shortcut for tf.add(a, b)
add_and_triple = adder_node * 3.
sess = tf.Session()
print(sess.run(add_and_triple, {a: 3, b: 4.5}))
```

```
//print result:
22.5
```



Device Placement

EXEMPLO 1: -- variable

```
with tf.device("/gpu:1"):
    v = tf.get_variable("v", [1])
```

"/cpu:0": The CPU of your machine.

"/gpu:0": The GPU of your machine, if you have one.


"/gpu:1": The second GPU of your machine, etc.

EXEMPLO 2: -- computation

```
with tf.device("/device:CPU:0"):
    # Operations created in this context will be pinned to the CPU.
    img = tf.decode_jpeg(tf.read_file("img.jpg"))

with tf.device("/device:GPU:0"):
    # Operations created in this context will be pinned to the GPU.
    result = tf.matmul(weights, img)
```

Device Placement





Code 3: Addition of tf.device

```
import tensorflow as tf

with tf.device('/cpu:0'):
    x = tf.placeholder(tf.float32, [None, 784])
with tf.device('/gpu:0'):
    W = tf.Variable(tf.zeros([784, 10]))
with tf.device('/gpu:1'):
    b = tf.Variable(tf.zeros([10]))
with tf.device('/cpu:0'):
    y_1 = tf.matmul(x, W)
with tf.device('/cpu:0'):
    y_2 = tf.add(y_1, b)
with tf.device('/gpu:2'):
    y = tf.nn.softmax(y_2)
```

TensorFlow runtime


TensorFlow in one picture





tensorboard

TensorBoard


GRAPHS

 Fit to screen

 Download PNG

Run log 

(1)

Session runs (0) 

Upload

Choose File

☐ Trace inputs

Color ☒ Structure

☐ Device

☐ XLA Cluster

☐ Compute time

☐ Memory


☐ TPU Compatibility

colors


same substructure

unique substructure


Const




addition




mult





x





y











Conferences



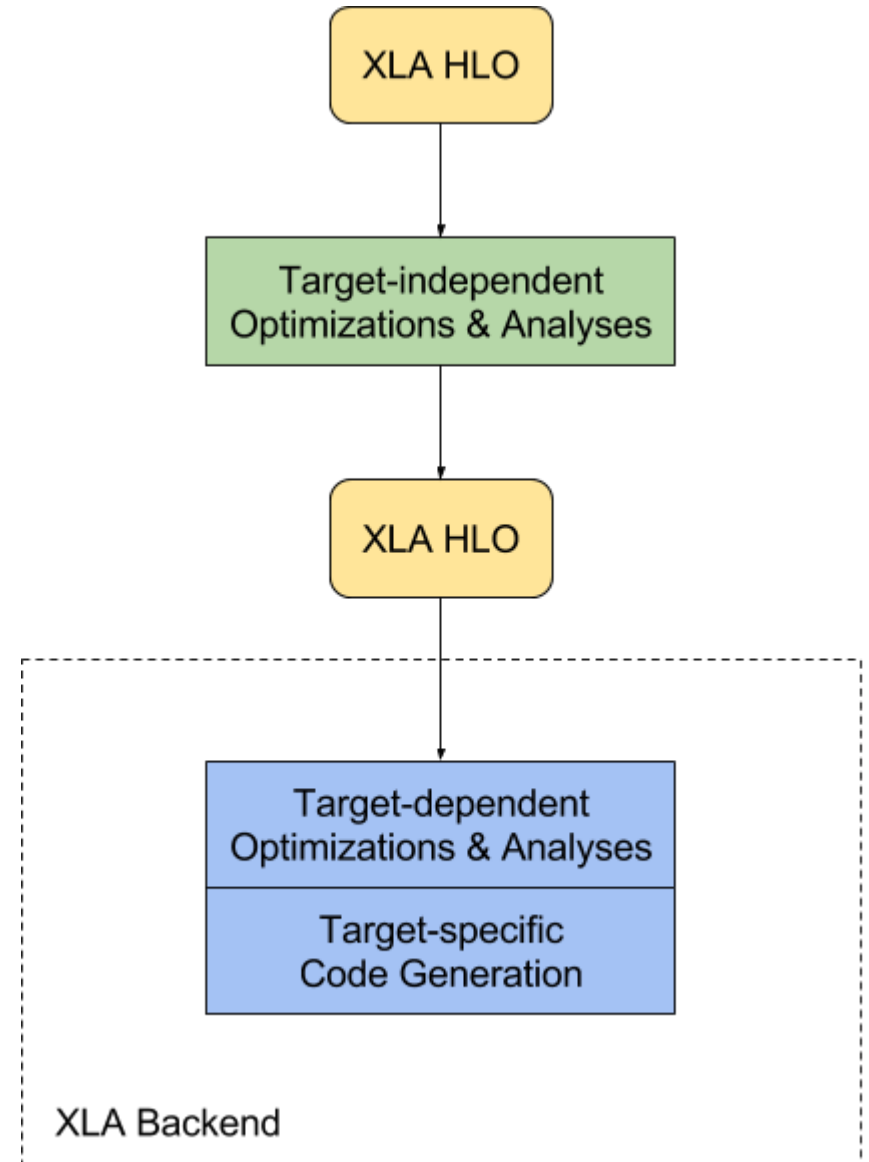
XLA

XLA (Accelerated Linear Algebra) is a domain-specific compiler for linear algebra that optimizes TensorFlow computations.

***OPTS**

- Improve execution speed
- Improve memory usage
- Reduce reliance on custom Ops
- Reduce mobile footprint.
- Improve portability

<https://goo.gl/rAsyHB>



XLA

```
//tensorflow/compiler/xla
```

```
//tensorflow/compiler/tf2xla
```


```
//tensorflow/compiler/jit
```

```
//tensorflow/compiler/aot
```

//tensorflow/compiler/xla

In general, XLA module is responsible to convert a given XLA Graph into binary by first lowering it into an LLVM IR, and after compiling the LLVM IR to given binary of a specific architecture (e.g. x86).

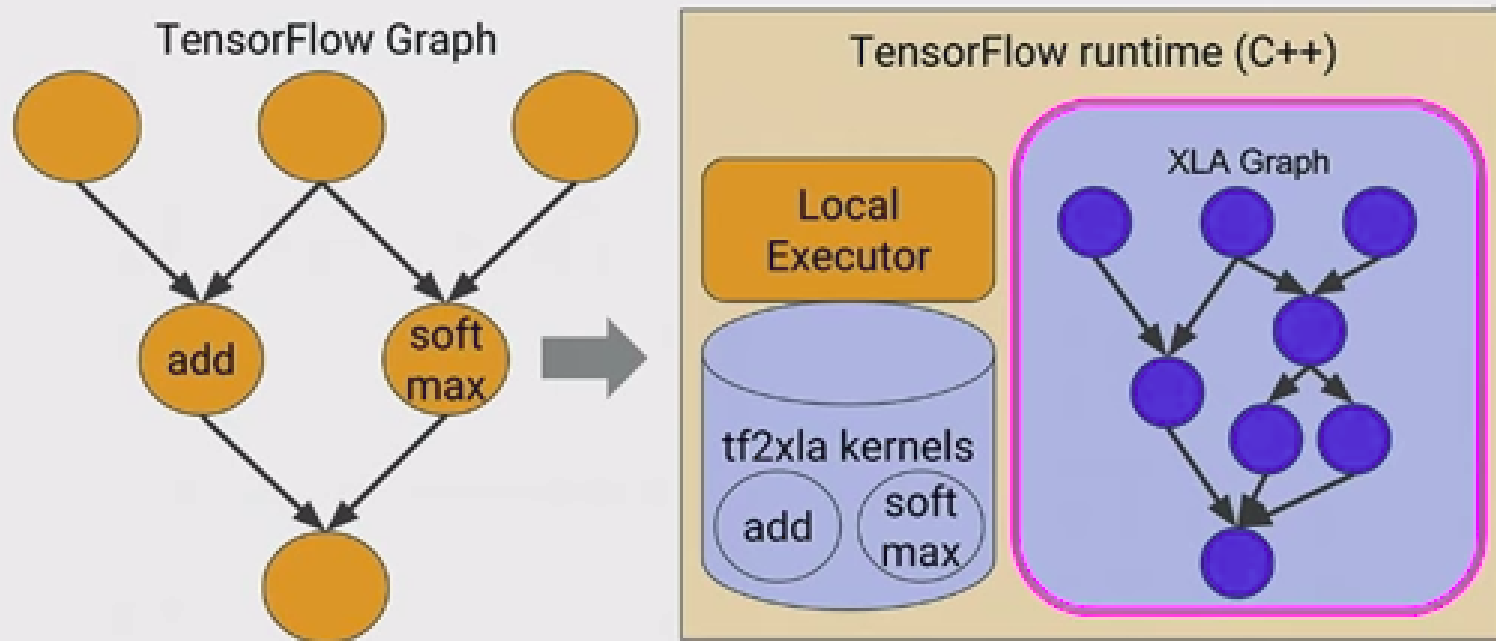
XLA in one picture



//tensorflow/compiler/tfa2xla

The tfa2xla module is responsible to convert an TFG into an XLA Graph. The example below convert the entire TFG into a XLA Graph (a common case when using AOT).

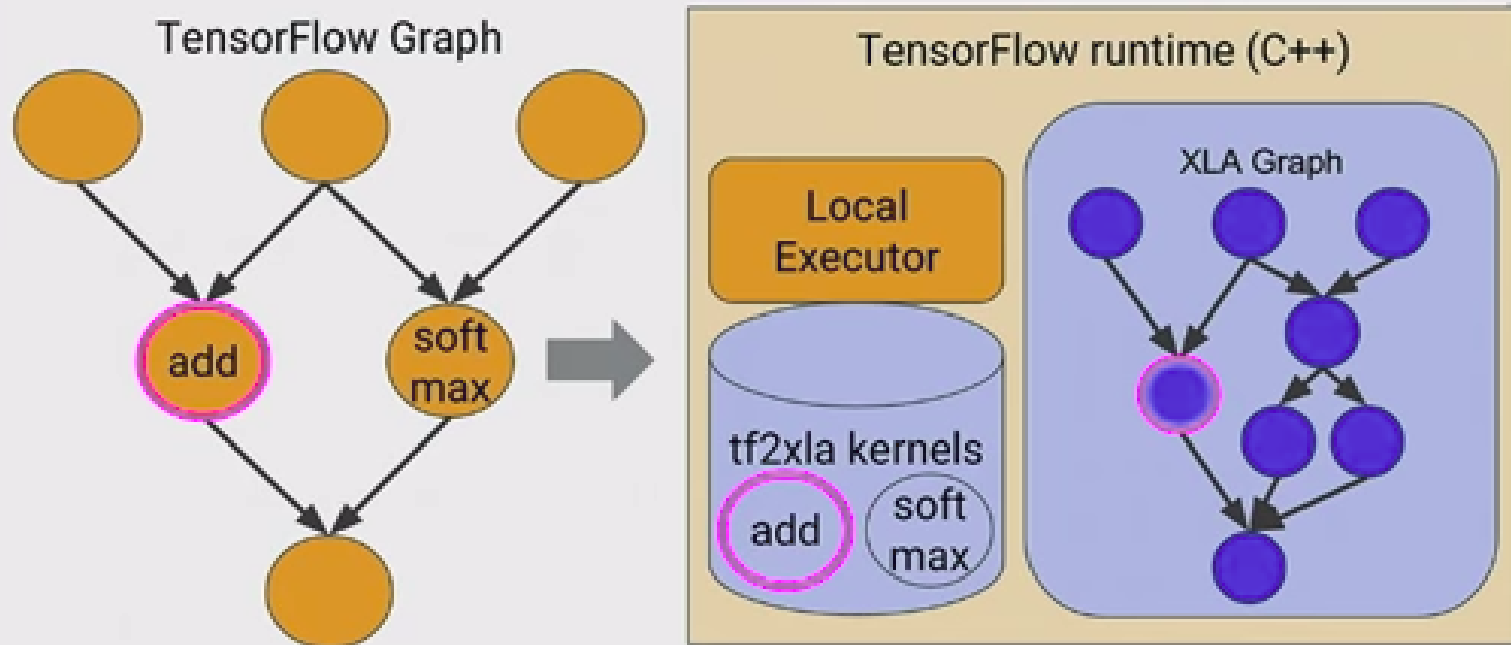
tf2xla: Symbolic graph execution



//tensorflow/compiler/tfa2xla

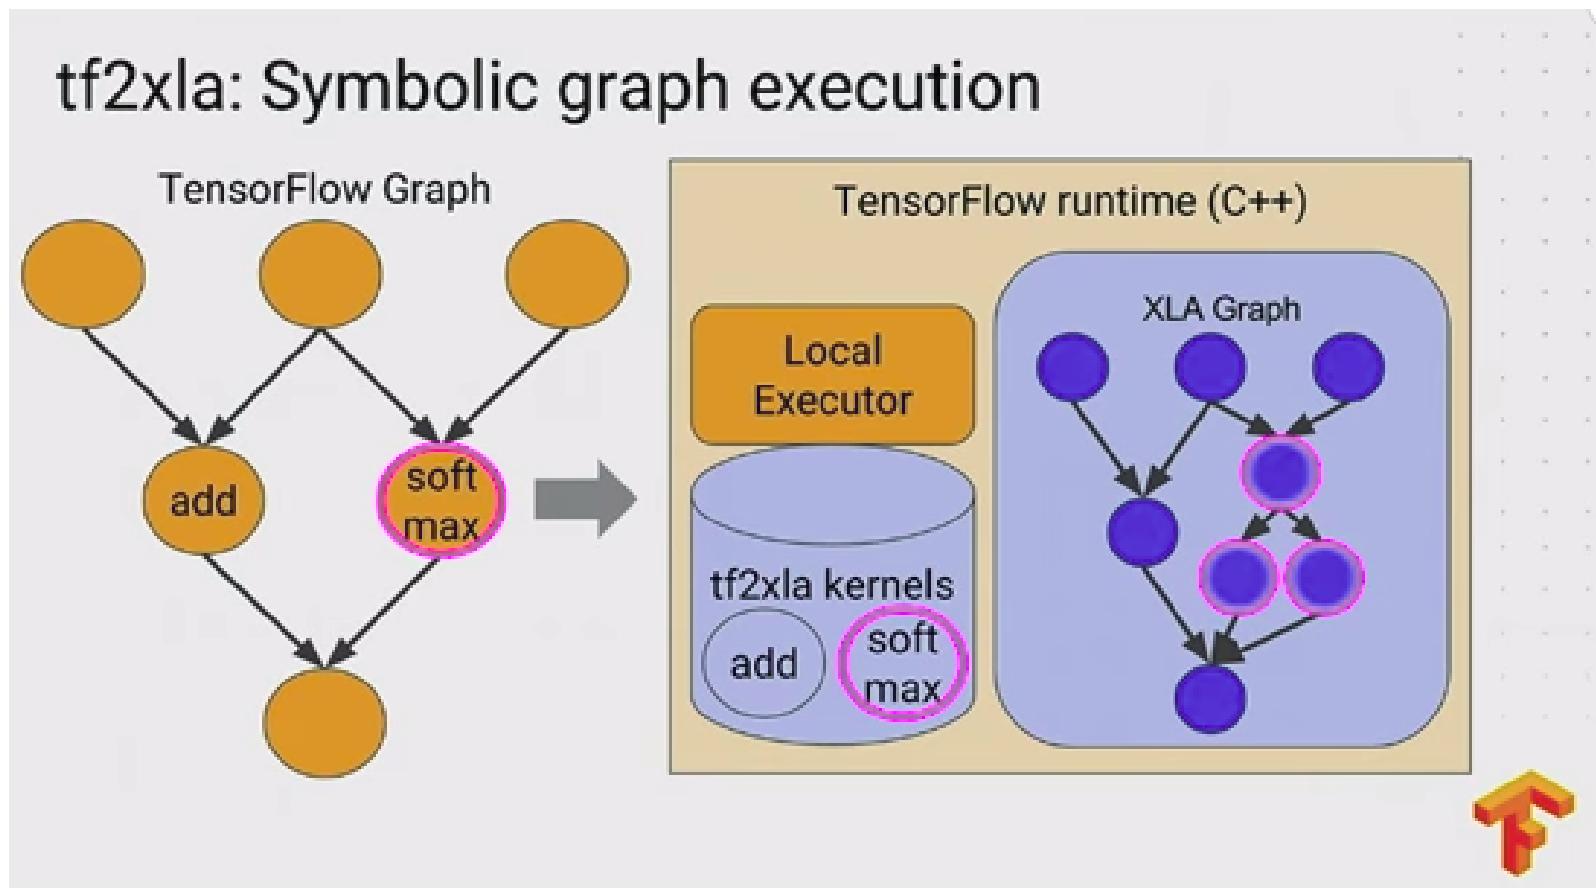
Each node of the TFG is converted into an XLA Graph's node. Observe that the TFG add node is converted into an XLA add node.

tf2xla: Symbolic graph execution




//tensorflow/compiler/tfa2xla

The TFG soft max node is converted into three XLA node.




//tensorflow/compiler/jit

JIT: Compile and run TF clusters



//tensorflow/compiler/jit

JIT: Multiple clusters



//tensorflow/compiler/jit

Turning on JIT compilation

<https://www.tensorflow.org/versions/master/experimental/xla/jit>

Whole session

```
config = tf.ConfigProto()  
config.graph_options.optimizer_options.global_jit_level = tf.OptimizerOptions.ON_1  
sess = tf.Session(config=config) # All supported ops compiled with XLA.
```


Manual scoped

```
jit_scope = tf.contrib.compiler.jit.experimental_jit_scope  
x = tf.placeholder(np.float32)  
with jit_scope():  
    y = tf.add(x, x) # The "add" op will be compiled with XLA.
```




//tensorflow/compiler/aot

Ahead-of-time compilation



//tensorflow/compiler/aot

tfcompile: Graph compiler



//tensorflow/compiler/aot

tfcompile

Write code to call the computation:

```
#include "myproject/tests/test_matmul.h" // generated header

int main(int argc, char** argv) {
    foo::TestMatMul matmul;

    // Set up args and run the computation.
    const float args[12] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12};
    std::copy(args + 0, args + 6, matmul.arg0_data());
    std::copy(args + 6, args + 12, matmul.arg1_data());
    matmul.Run();

    // Check results
    CHECK_EQ(matmul.result0(0, 0), 58);
    return 0;
}
```

//tensorflow/compiler/aot

Smaller binaries on mobile

Binary size reduction on android-arm (stacked LSTM, 3 deep, 60 wide)

Original: 2.6MB (1MB runtime + 1.6MB graph)

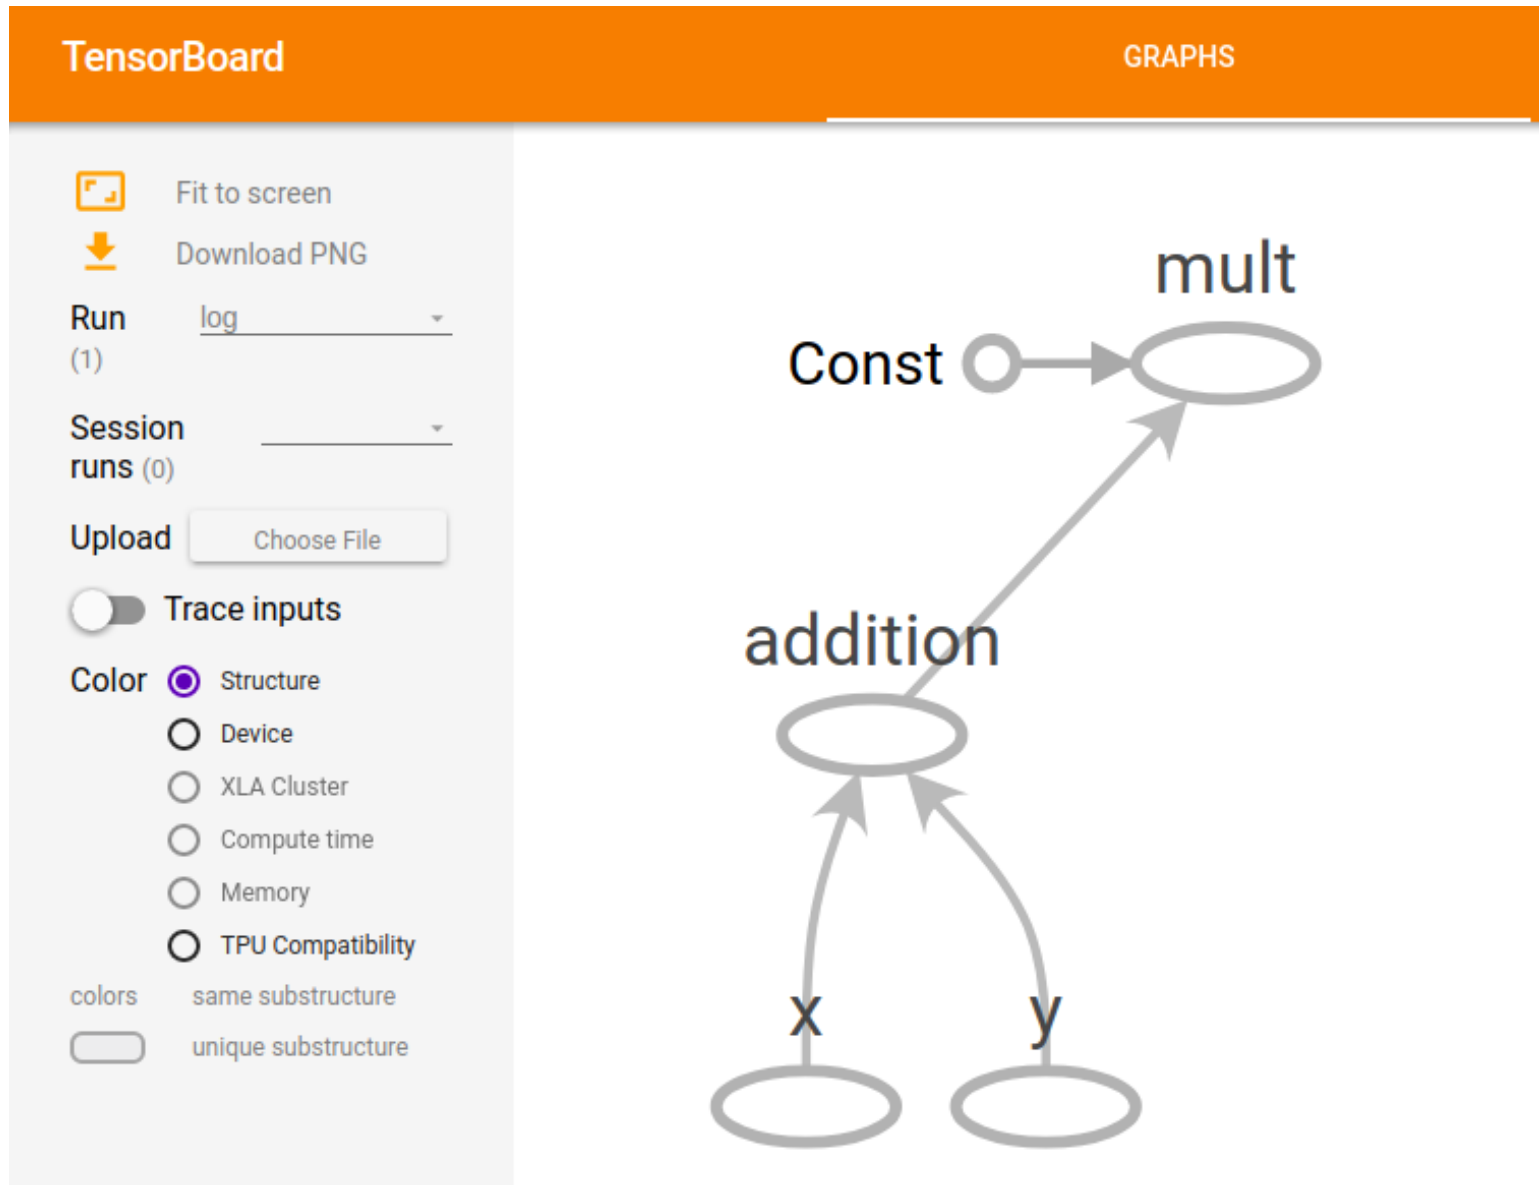
Compiled: **600KB** (**272KB code** + 330KB weights)

Example JIT

```
1 import tensorflow as tf
2
3 jit_scope = tf.contrib.compiler.jit.experimental_jit_scope
4
5 X = tf.placeholder(tf.float32, name="x")
6 Y = tf.placeholder(tf.float32, name="y")
7 value = tf.constant(5, tf.float32)
8
9 addition = tf.add(X, Y, name="addition")
10
11 with jit_scope():
12     mult = tf.multiply(addition, value, name="mult")
13
14 session = tf.Session()
15 result = session.run(mult, feed_dict={X: [5,2,1], Y: [10,6,1]})
16
17 # writting the computational graph for tensorboard
18 writer = tf.summary.FileWriter("/home/rafael/Samsung/tests/log/", session.graph)
19
20 print(result)
```

Print --> [75. 40. 10.]

tensorboard




Example of Protobuf

```
node {  
  name: "addition"  
  op: "Add"  
  input: "X"  
  input: "Y"  
  device: "CPU:0"  
  attr {  
    key: "T"  
    value {  
      type: DT_FLOAT  
    }  
  }  
}
```


Execution -- *python example.py*

- `Y = _Arg[T=DT_FLOAT, index=1, CPU:0"]()` is dead: 0




Execution -- *python example.py*

- `X = _Arg[T=DT_FLOAT, index=0, CPU:0"]()` is dead: 0




Execution -- *python example.py*

- `Const = Const[dtype=DT_FLOAT, value=Tensor<type: float shape: [] values: 5>, CPU:0"]()` is dead: 0



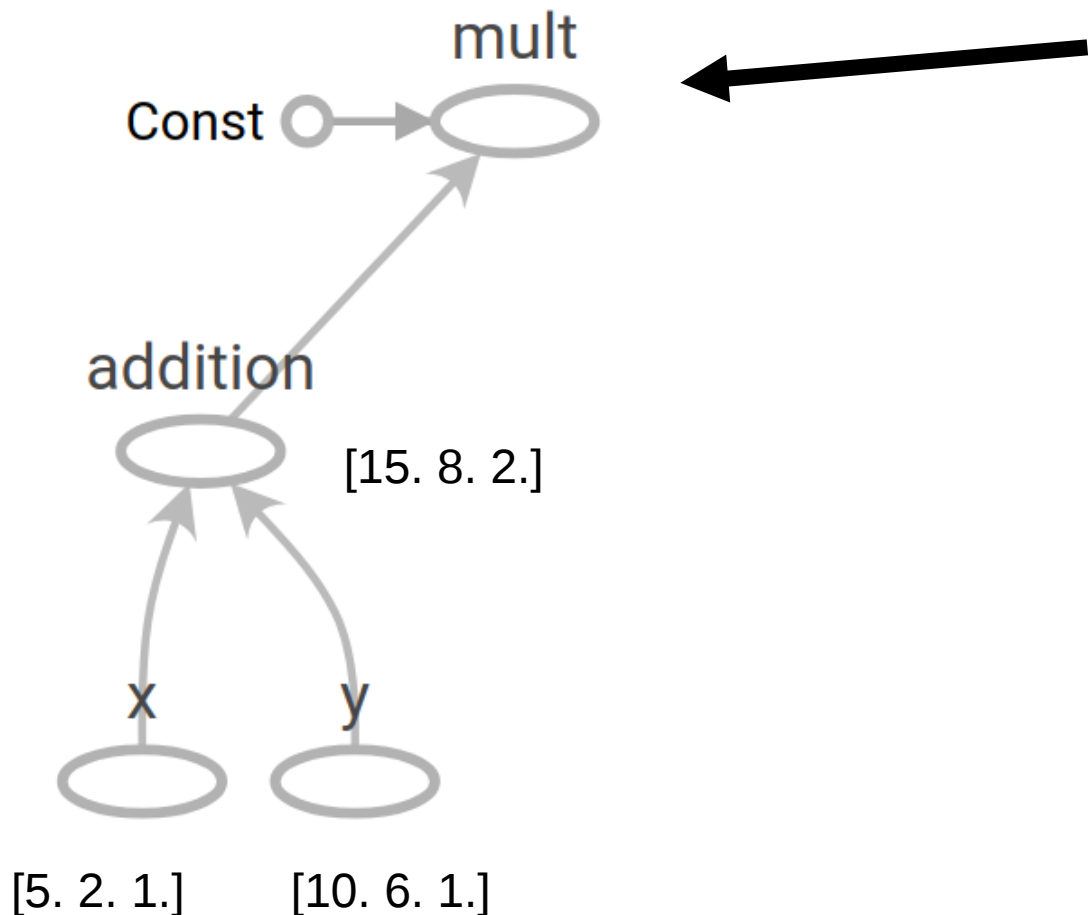
Execution -- *python example.py*

- `addition = Add[T=DT_FLOAT, CPU:0"](X,Y)` is dead: 0



Execution -- *python example.py*

- **[XLA]** `mult = multiply[Targs=[DT_FLOAT, DT_FLOAT],
Tresults=[DT_FLOAT], function=funcMult, _device="CPU:0"]`(`addition`,
`Const`) is dead: 0



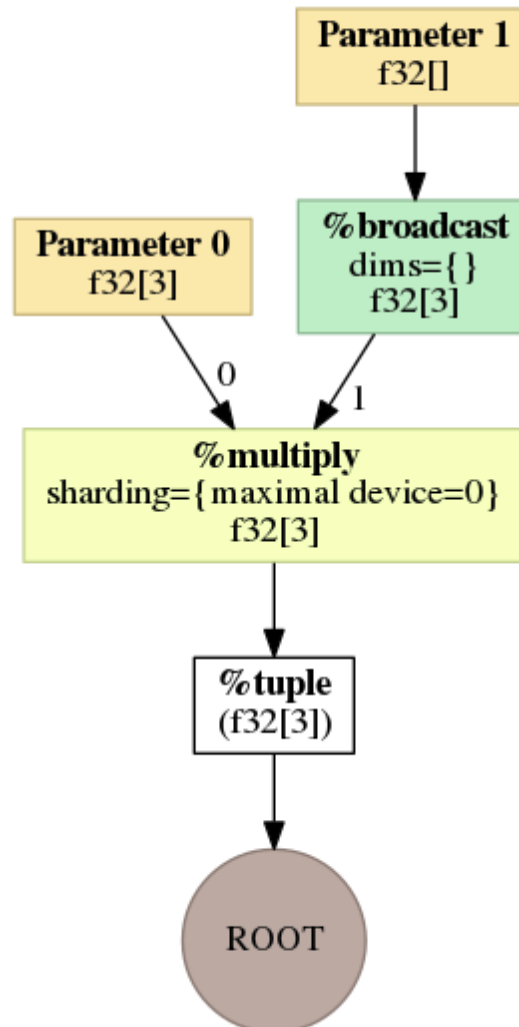
XLA compilation

- **XlaCompilationCache::Compile XLA JIT compilation cache**
 - Signature: funcMult ,float[3],float[]
 - num_inputs = 2
 - 0: dtype=float present=1 shape=[3]
 - 1: dtype=float present=1 shape=[]
 - num_outputs = 1
 - 0: dtype=1

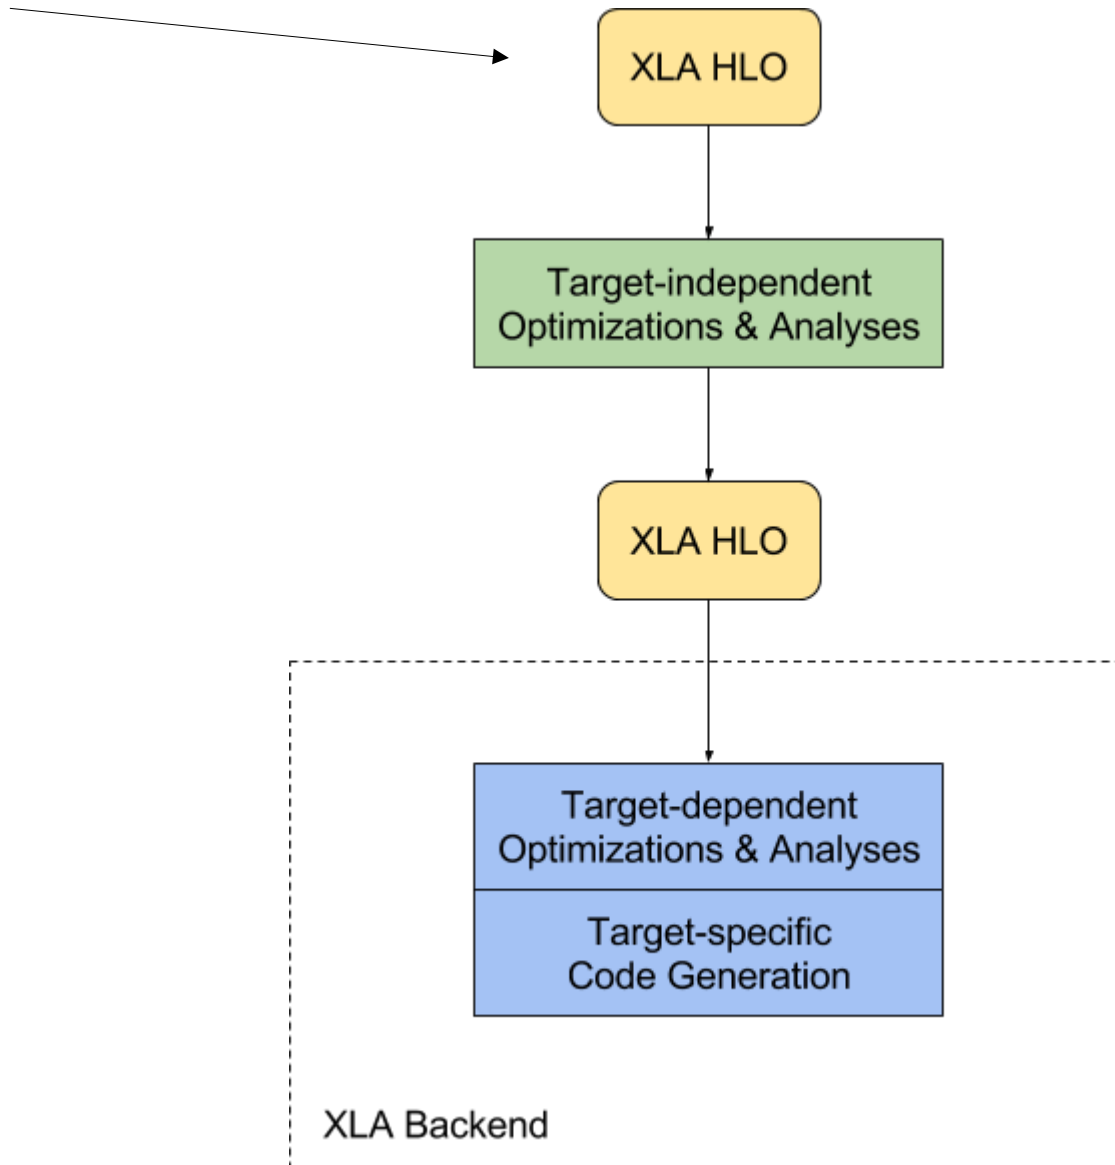
Compilation cache miss for the signature above!!

XLA Graph

XLA Graph -- Computation funcMult



XLA compilation



HLO IR

HloModule funcMult:

ENTRY

%cluster_0[_XlaCompiledKernel=true,_XlaNumConstantArgs=0,_XlaNumResourceArgs=0].v4

(arg0: f32[3], arg1: f32[]) -> (f32[3]) {

 %arg0 = f32[3]{0} parameter(0)

 %arg1 = f32[] parameter(1)

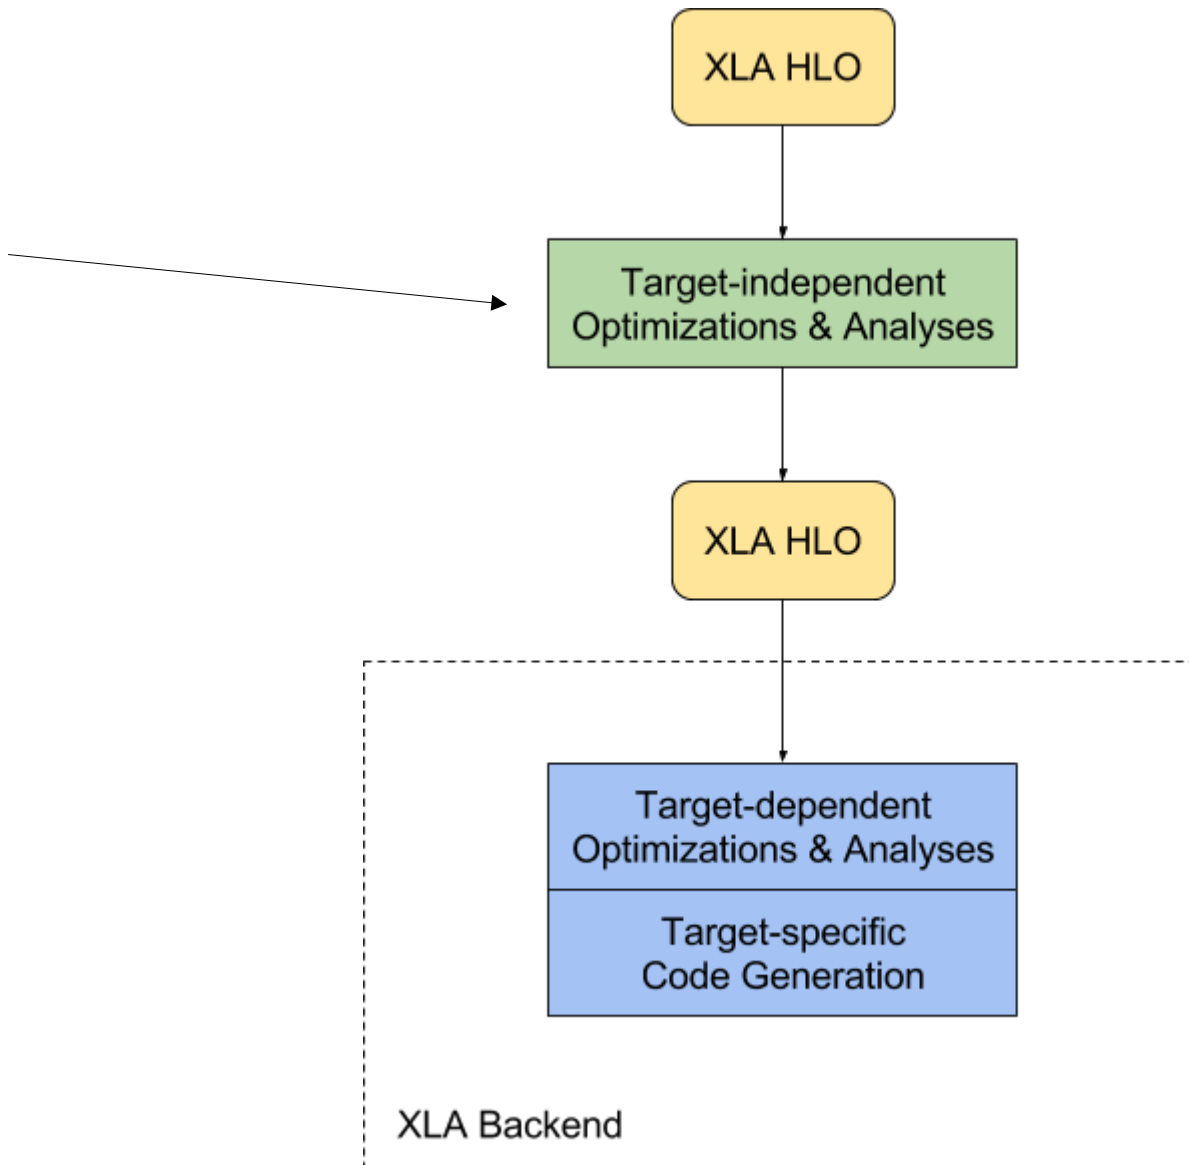
 %broadcast = f32[3]{0} broadcast(f32[] %arg1), dimensions={}

 %multiply = f32[3]{0} multiply(f32[3]{0} %arg0, f32[3]{0} %broadcast), sharding={maximal
device=0} # metadata=op_type: "Mul" op_name: "mult"

 ROOT %tuple = (f32[3]{0}) tuple(f32[3]{0} %multiply)

}

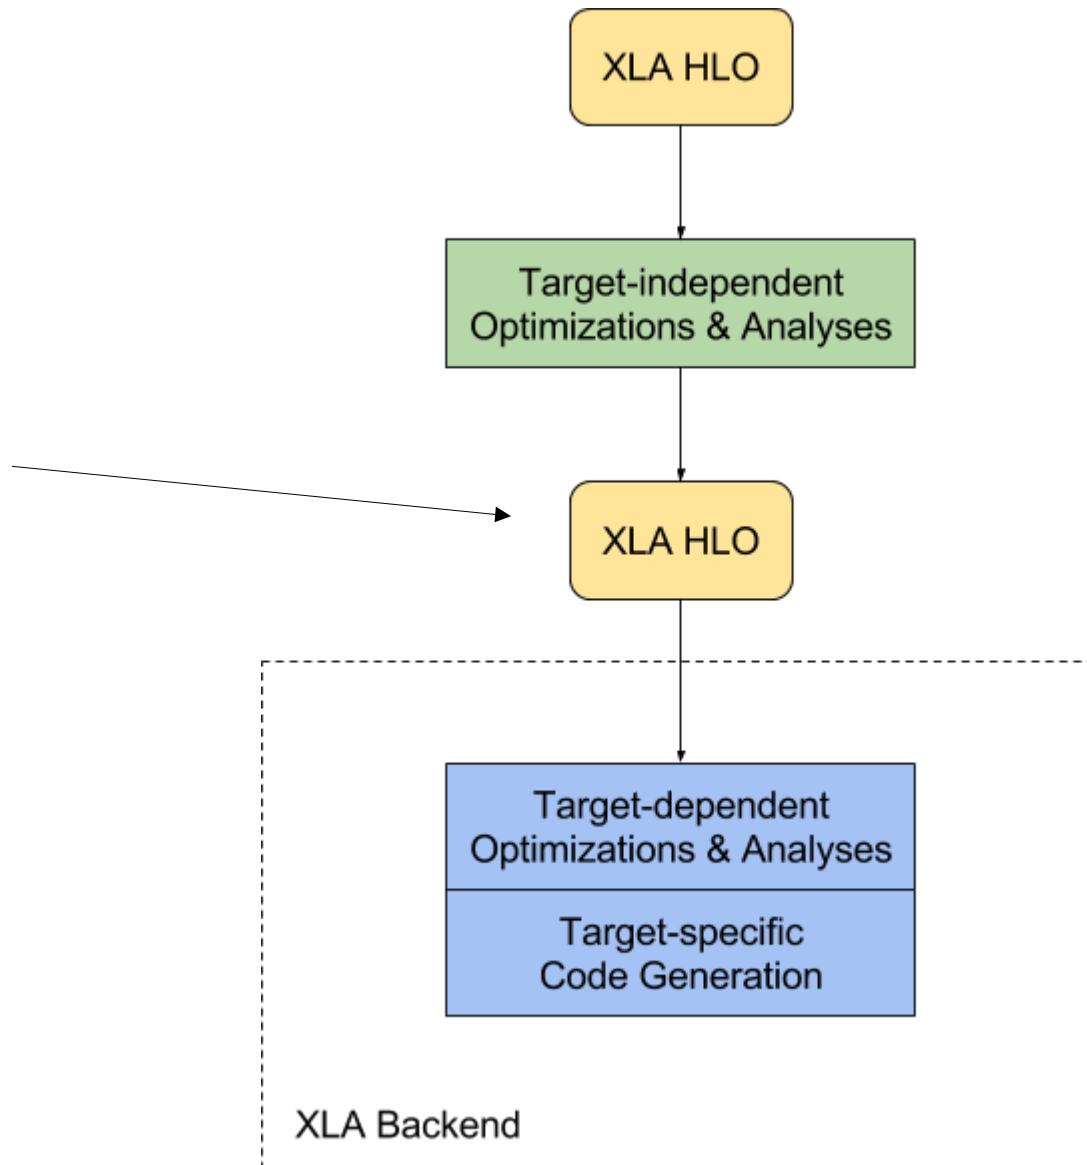
XLA compilation



XLA compilation -- OPT

- HLO pass pipeline CPU
- HLO pass CallInliner
- HLO pass convolution-canonicalization
- HLO pass simplification
- HLO pass batchnorm_rewriter
- HLO pass algsimp
- HLO pass tuple-simplifier
- HLO pass dce
- HLO pass reshape-mover
- HLO pass constant_folding
- HLO pass transpose-folding
- HLO pass cse
- HLO pass fusion
- HLO pass layout-assignment
- HLO pass algsimp
- HLO pass cse
- HLO pass cpu-parallel-task-assigner
- HLO pass dce
- HLO pass copy-insertion
- HLO pass dce
- HLO pass flatten-call-graph

XLA compilation



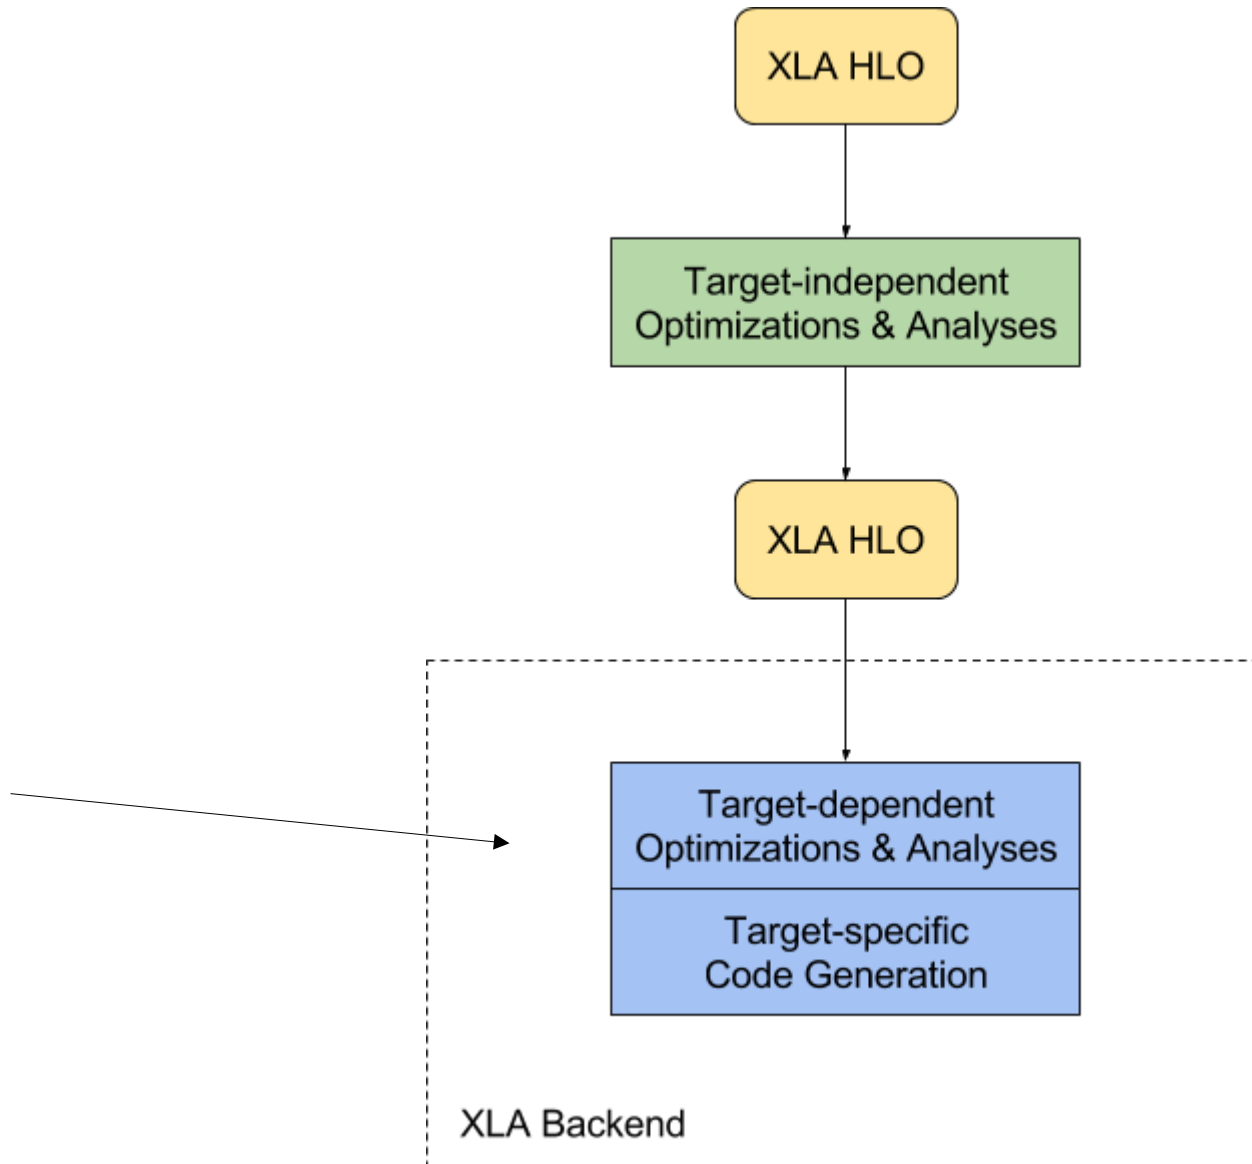
HLO IR

HloModule funcMult:

ENTRY

```
%cluster_0[_XlaCompiledKernel=true,_XlaNumConstantArgs=0,_XlaNumResourceArgs=0].v4 (arg0: f32[3], arg1: f32[]) -> (f32[3]) {  
  %arg0 = f32[3]{0} parameter(0)  
  %arg1 = f32[] parameter(1)  
  %fusion = f32[3]{0} fusion:kLoop(f32[3]{0} %arg0, f32[] %arg1), calls=  
%fused_computation # metadata=op_type: "Mul" op_name: "mult"  
  %fused_computation (arg0.param_0: f32[3], arg1.param_1: f32[]) -> f32[3] {  
    %arg0.param_0 = f32[3]{0} parameter(0)  
    %arg1.param_1 = f32[] parameter(1)  
    %broadcast.1 = f32[3]{no layout} broadcast(f32[] %arg1.param_1), dimensions={}  
    ROOT %multiply.1 = f32[3]{0} multiply(f32[3]{0} %arg0.param_0, f32[3]{no layout}  
%broadcast.1), sharding={maximal device=0} # metadata=op_type: "Mul" op_name:  
"mult"  
  }  
  ROOT %tuple = (f32[3]{0}) tuple(f32[3]{0} %fusion)  
}
```

XLA compilation



XLA compilation --LLVM IR

```
; ModuleID = '__compute_module'
source_filename = "__compute_module"
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"
```

```
define void @multFunc(i8* align 8 dereferenceable(8) %retval, i8* noalias %run_options, i8** noalias %params, i8** noalias %temps, i64* noalias %prof_counters) #0 {
```

```
entry:
```

```
    %fusion.invar_address.dim.0 = alloca i64
    %0 = getelementptr inbounds i8*, i8** %params, i64 0
    %arg0.untyped = load i8*, i8** %0, !invariant.load !0, !dereferenceable !1, !align !2
    %1 = bitcast i8* %arg0.untyped to [3 x float]*
    %2 = getelementptr inbounds i8*, i8** %params, i64 1
    %arg1.untyped = load i8*, i8** %2, !invariant.load !0, !dereferenceable !3, !align !2
    %3 = bitcast i8* %arg1.untyped to float*
    %4 = getelementptr inbounds i8*, i8** %temps, i64 0
    %5 = load i8*, i8** %4, !invariant.load !0, !dereferenceable !1, !align !2
    %fusion = bitcast i8* %5 to [3 x float]*
    store i64 0, i64* %fusion.invar_address.dim.0
    br label %fusion.loop_header.dim.0
```

```
fusion.loop_header.dim.0:                ; preds = %fusion.loop_body.dim.0, %entry
    %fusion.indvar.dim.0 = load i64, i64* %fusion.invar_address.dim.0
    %6 = icmp uge i64 %fusion.indvar.dim.0, 3
    br i1 %6, label %fusion.loop_exit.dim.0, label %fusion.loop_body.dim.0
```

```
fusion.loop_body.dim.0:                  ; preds = %fusion.loop_header.dim.0
    %7 = getelementptr inbounds [3 x float], [3 x float]* %1, i64 0, i64 %fusion.indvar.dim.0
    %8 = load float, float* %7, !invariant.load !0, !noalias !4
    %9 = load float, float* %3, !invariant.load !0, !noalias !4
    %10 = fmul fast float %8, %9
    %11 = getelementptr inbounds [3 x float], [3 x float]* %fusion, i64 0, i64 %fusion.indvar.dim.0
    store float %10, float* %11, !alias.scope !4, !noalias !7
    %invar.inc = add nuw nsw i64 %fusion.indvar.dim.0, 1
    store i64 %invar.inc, i64* %fusion.invar_address.dim.0
    br label %fusion.loop_header.dim.0
```

```
fusion.loop_exit.dim.0:                  ; preds = %fusion.loop_header.dim.0
    %tuple = bitcast i8* %retval to [1 x i8]*
    %12 = getelementptr inbounds [1 x i8], [1 x i8]* %tuple, i64 0, i64 0
    %13 = bitcast [3 x float]* %fusion to i8*
    store i8* %13, i8** %12, !alias.scope !7, !noalias !4
    %prof_counter.computation = getelementptr i64, i64* %prof_counters, i64 0
    ret void
```

```
}
```



XLA compilation -- x86_64

0x00000000	movrax, qword ptr [rdx]
0x00000003	movrdx, qword ptr [rdx + 8]
0x00000007	movrcx, qword ptr [rcx]
0x0000000a	vmovss xmm0, dword ptr [rdx]
0x0000000e	vmulss xmm1, xmm0, dword ptr [rax]
0x00000012	vmovss dword ptr [rcx], xmm1
0x00000016	vmulss xmm1, xmm0, dword ptr [rax + 4]
0x0000001b	vmovss dword ptr [rcx + 4], xmm1
0x00000020	vmulss xmm0, xmm0, dword ptr [rax + 8]
0x00000025	vmovss dword ptr [rcx + 8], xmm0
0x0000002a	movqword ptr [rdi], rcx
0x0000002d	ret

Execution -- *python example.py*

- **[XLA]** `mult = multiply[Targs=[DT_FLOAT, DT_FLOAT],
Tresults=[DT_FLOAT], function=funcMult, _device="CPU:0"]`(`addition`,
`Const`) is dead: 0


funcMult found



Execution -- *python example.py*

- `retval(mult)`

[75. 40. 10.]




Example AOT

```
1 import tensorflow as tf
2
3 x_hold = tf.placeholder(tf.float32, name="x_hold")
4 y_hold = tf.placeholder(tf.float32, name="y_hold")
5 x_y_prod = tf.multiply(x_hold, y_hold, name="x_y_prod")
6
7 session = tf.Session()
```

TensorFlow Graph

It is necessary
to create the
protobuf from
the TensorFlow
Graph

tfmatmul.pb



Feeds and Fetches

```
feed {  
  id { node_name: "x_hold" }  
  shape {  
    dim { size: 2 }  
    dim { size: 3 }  
  }  
}
```

tfmatmul.config.pbtxt

```
feed {  
  id { node_name: "y_hold" }  
  shape {  
    dim { size: 3 }  
    dim { size: 2 }  
  }  
}
```

```
fetch {  
  id { node_name: "x_y_prod" }  
}
```

Invoke tfcompile

```
load("//third_party/tensorflow/compiler/aot:tfcompile.bzl", "tf_library")
```

```
# Use the tf_library macro to compile your graph into executable  
code.
```

```
tf_library(  
    name = "tfmatmul",  
    cpp_class = "foo::bar::MatMulComp",  
    graph = "tfmatmul.pb",  
    config = "tfmatmul.config.pbtxt",  
)
```

Invoke the tfcompile passing this script as parameter

tfcompile OPT

tfcompile makes use of XLA optimizations by converting the protobuf into HLO IR.

tfcompile results (header and binary)


```
namespace foo {
namespace bar {

class MatMulComp {
public:
    enum class AllocMode {
        ARGS_RESULTS_AND_TEMPS, // Allocate arg, result and temp
        RESULTS_AND_TEMPS_ONLY, // Only allocate result and temp
    };

    MatMulComp(AllocMode mode = AllocMode::ARGS_RESULTS_AND_TEMPS);
    ~MatMulComp();

    bool Run();
    void** args();
    void set_arg0_data(float* data);
    float* arg0_data();
    float& arg0(size_t dim0, size_t dim1);
    void set_arg1_data(float* data);
    float* arg1_data();
    float& arg1(size_t dim0, size_t dim1);
    void** results();
    float* result0_data();
    float& result0(size_t dim0, size_t dim1);
};

} // end namespace bar
} // end namespace foo
```



tfcompile has as
result one
header file and
one binary (e.g.
x86).

Using the header and binary files

```
#define EIGEN_USE_THREADS
#define EIGEN_USE_CUSTOM_THREAD_POOL

#include <iostream>
#include "third_party/eigen3/unsupported/Eigen/CXX11/Tensor"
#include "tensorflow/compiler/aot/tests/test_graph_tfmatmul.h" //
generated

int main(int argc, char** argv) {
    Eigen::ThreadPool tp(2); // Size the thread pool as appropriate.
    Eigen::ThreadPoolDevice device(&tp, tp.NumThreads());

    foo::bar::MatMulComp matmul;
    matmul.set_thread_pool(&device);

    // Set up args and run the computation.
    const float args[12] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12};
    std::copy(args + 0, args + 6, matmul.arg0_data());
    std::copy(args + 6, args + 12, matmul.arg1_data());
    matmul.Run();

    // Check result
    if (matmul.result0(0, 0) == 58) {
        std::cout << "Success" << std::endl;
    } else {
        std::cout << "Failed. Expected value 58 at 0,0. Got:"
                    << matmul.result0(0, 0) << std::endl;
    }

    return 0;
}
```

Compile the code linking the binary generated by tfcompile and execute.